

The Enterprise Guide to Agentic Development

The operating model engineering organizations use to scale AI output while preserving quality, governance and senior engineering judgment. Applied by engineering teams across Fortune 100 companies, DAX 40 enterprises and high-growth technology organizations in regulated and security-sensitive environments.



Executive Summary

AI coding tools are rapidly increasing engineering output across organizations. Most teams, however, adopt these tools without establishing the governance, verification standards and responsibility models required to use them safely at scale.

Output is increasing faster than engineering control.

Organizations implementing structured AI-assisted development typically achieve:

- 40% engineering velocity improvement within the first month
- Reduction in PR cycle time and review overhead
- Team-wide AI adoption increasing from approximately 50% to over 95%
- Improved code quality and reduced AI-introduced technical debt
- Faster onboarding and knowledge transfer across engineering teams

Without a standardized operating model, most organizations capture only a fraction of AI productivity potential while increasing quality, security, and architectural risk.

High-performing engineering organizations address this by explicitly defining how responsibility is distributed between engineers and AI systems. They formalize when work should be delegated to AI, when outputs must be reviewed and where human judgment remains mandatory.

This guide introduces the operating model leading engineering teams use to safely scale AI-assisted development across teams, codebases and entire engineering organizations.

Most organizations deploying AI coding tools increase output immediately. What they do not increase is engineering control, verification standards, or governance. As a result, junior developers generate code they cannot reliably validate, while senior engineers see limited productivity gains. Organizations lack a shared operating model defining what AI should execute and what must remain under human judgment.

The performance gap is measurable. Research¹ shows senior engineers improve productivity by roughly 22% when using AI coding tools, while junior developers see improvements closer to 4%. The tools are identical. The difference is methodology and responsibility allocation across the team.

AI amplifies existing engineering capability. Strong engineers become significantly more effective. Developing engineers gain limited advantage without structure. The challenge is not tool adoption, but systematically elevating team-wide engineering performance.

The Cost of Getting it Wrong

Quality Degradation:

Junior developers accept AI suggestions without understanding the implications. Security vulnerabilities, performance issues, and architectural violations slip through. The codebase accumulates technical debt faster than before AI adoption.

Inconsistent Practices:

Some team members use AI extensively. Others resist entirely. There is no shared language for AI-assisted work and no agreed standards for when delegation is appropriate.

Missed Competitive Advantage:

Your competitors are developing systematic approaches to AI-assisted development. Without methodology, you capture roughly 10% of the potential productivity gain while assuming 100% of the risk.

Senior Engineer Frustration:

Your most experienced engineers see AI tools as either a threat to be resisted or a toy that does not match their workflow. Neither perspective captures the actual opportunity.

¹ <https://jellyfish.co/research/ai-coding/>

The Framework: Delegate, Review, Own

High-performing engineering organizations do not “use AI.” They define how responsibility is distributed between engineers and AI systems. This framework is that operating model. Every task in agentic development falls into one of three modes.

Delegate

- The AI executes. You provide a clear specification and constraints. Senior engineers delegate better because they remove ambiguity upfront.
- Use for repeatable work where correctness can be verified quickly: Boilerplate, test scaffolding, migrations, documentation, log analysis, standard implementations.
- Success metric: clear spec, small diff, predictable output, minimal back-and-forth.

Review

- The AI proposes. You verify. You are the quality gate. Senior engineers review faster because they recognize correct patterns and failure modes early.
- Use when risk is non-trivial: Edge cases, security implications, architectural fit, failure handling, meaningful test coverage.
- Review focuses on failure modes: security, data handling, edge cases, and unintended scope.

Own

- The AI cannot make this decision. Human judgment is the value.
- Use for decisions that set direction and standards: Product direction, architecture trade-offs, risk acceptance, team conventions, engineering standards for AI usage.
- Own defines the rules: what AI may touch, how it must be verified, and what requires approval.

Applying the Framework Across Development Phases

Phase	Delegate	Review	Own
Plan	Draft specifications, map dependencies	Validate scope, verify estimates	Vision, trade-offs, approval
Design	Generate component skeletons	Check accessibility, performance	System design, architecture
Build	Infrastructure code, configurations	Audit security, cost implications	Architecture decisions
Test	Generate test suites	Reject shallow coverage	Test strategy, data approach
Review	Automated style and pattern checks	Verify intent, audit for drift	Standards, risk acceptance
Document	Generate changelogs, technical docs	Verify accuracy	Documentation structure
Deploy	Execute deployments, monitor logs	Audit configurations	Release decisions, incident response

If a team disagrees on a cell, that cell belongs in Own until standards are defined.

Why Self-Learning Is Not Sufficient

Engineering teams can learn AI tools independently. They watch tutorials, experiment, and develop personal workflows. That does not scale.

The failure mode is inconsistency. Each developer invents their own approach, standards drift across the codebase, and best practices discovered by one engineer do not transfer to others. The organization captures a fraction of the potential value and absorbs most of the risk.

Structured training provides four outcomes that self-learning rarely produces:

I. Senior Engineer Buy-In

Skeptical senior engineers become actual advocates when they see the methodology protects judgment rather than replacing it. This eliminates the cultural split between “AI power users” and “AI resisters.”

III. Reusable Playbooks

Teams leave with concrete artifacts: verification checklists, prompt planning templates, documentation standards, and conventions for what AI may change and how it must be reviewed.

II. Shared Operating Model

Teams adopt a common language for responsibility allocation. Our proposed “Delegate”, “Review”, and “Own” become enforceable defaults instead of personal preference.

IV. Faster Organizational Adoption

What takes months of individual trial and error becomes weeks of consistent rollout. Output increases without sacrificing quality because verification and ownership boundaries are defined upfront.

Six Actions High-Performing Teams Take

I. Establish Documentation for AI

Create documentation that helps AI tools understand your codebase. This includes project conventions, testing approaches, common patterns, and known quirks. When AI tools have context, they produce more accurate output

II. Write Implementation Plans Before Prompting

Plan before you prompt. Outline the goal, break it into discrete tasks, and execute one task at a time. This prevents scope creep and produces cleaner implementations.

III. Reset Early When AI Output Goes Wrong

When AI produces incorrect output, do not stack corrections. Discard the attempt and restart with a clearer specification. Iterating on a bad foundation creates fragile code.

IV. Decompose Work Into Small Units

AI performs best on focused, well-defined tasks. Asking for entire features produces sprawling output. Asking for specific components produces accurate implementations.

V. Create Verification Checklists for Junior Developers

Junior developers cannot reliably verify AI output by instinct. Provide explicit checks: no hardcoded credentials, error states handled, failure cases tested, permissions follow least-privilege.

VI. Protect Core Engineering Skills

Do not delegate the skills that create leverage: problem decomposition, debugging, system design, trade-off analysis. Let AI handle execution while engineers maintain judgment.

What Happens When Teams Apply This

Across Fortune 100 companies, DAX 40 enterprises, and high-growth technology organizations, the pattern is consistent. When teams adopt a shared operating model for Delegate, Review, and Own, three things happen quickly: output increases, quality stabilizes, and senior engineers stop resisting.

- 40% velocity improvement within the first month
- Team-wide AI adoption increasing from approximately 50% to over 95%
- Cleaner codebases with enforced standards
- Fewer review cycles and reduced back-and-forth
- Faster onboarding for new team members
- Senior engineers moving from skepticism to advocacy

What changes in practice:

- PRs get smaller and easier to review
- Teams converge on shared prompting and verification standards
- Senior engineers spend less time fixing AI mistakes and more time owning architecture

Representative implementation snapshots:

- A regulated enterprise engineering team reduced average PR cycle time from 4 days to 1.5 days after standardizing review and verification practices.
- A high-growth SaaS organization shipped a major feature in 3 weeks that had previously been estimated at 8 weeks by moving work into smaller, verifiable units.
- A large multi-team engineering org standardized AI development practices across 200+ engineers in under a month by rolling out shared playbooks and checklists.

What Training Covers

This white paper explains the operating model. Training turns it into a repeatable capability inside your organization. We align on standards, apply the framework to real work, and leave your team with internal playbooks that continue to work after the workshop ends.

I. Tool Proficiency

We train your team on the tools you use or plan to adopt, including Cursor (Tab, Composer, multi-file context, MCP servers) and agentic workflows (Claude Code, Codex). The focus is practical execution in real codebases, not feature tours.

II. Infrastructure Development

We help teams set up the foundation that makes AI effective at scale: shared documentation, reusable templates, and integrations with existing systems. The goal is repeatable workflows across the organization, not individual hacks.

III. Judgment Development

When to delegate, when to review, and what must be owned. How to verify AI output quickly and reliably. How to develop junior engineers without creating dependency.

IV. Application to Your Codebase

We work on your actual code under NDA. Teams apply the framework to real work, align on conventions, and leave with internal playbooks and standards. For regulated environments, we incorporate the vital compliance constraints that matter.

V. Cross-Functional Coverage

AI-assisted development is a team capability. We include developers, QA, DevOps, and engineering leadership so everyone shares the same vocabulary, standards, and review expectations.

VI. Security & Compliance

AI-assisted development in regulated environments. Code auditing practices. Ensuring AI-generated code meets security standards. Managing sensitive data and credentials in AI workflows.

What Training Does Not Cover

We qualify teams upfront. This training is for organizations ready to operationalize AI-assisted development, not explore it. The goal is implementation that sticks: standards, workflows, and internal capability your team can run without us.

I. Basic ‘What is AI’ Introductions

We assume your team is ready to implement, not evaluate. If your organization is still deciding whether AI coding tools are worth adopting, that decision should happen before we engage.

II. Generic AI/ML Theory

We focus on practical application, not academic foundations. Your engineers do not need to understand transformer architectures to use AI coding tools effectively.

III. Building AI Models From Scratch

We teach teams to use agentic tools safely and effectively, not to become ML engineers. The goal is maximum leverage from existing tools, not building new models.

IV. One-Size-Fits-All Templates

Every engagement is customized to your stack, your team, and your constraints. Generic playbooks fail because they ignore the context that makes your organization unique.

V. Hype or Speculation

We deliver frameworks you can apply immediately, not predictions about AGI or the future of work. Everything we teach is grounded in practices that work today.

VI. Consulting Dependencies

We do not deliver slide decks while your team watches. Sessions are hands-on, with engineers applying concepts to real code in real time. The goal is complete internal capability, not external dependency.

Training Formats

On-Site

We come to your office, anywhere in the world. Minimum two trainers for every engagement. Custom curriculum based on your technology stack, team composition, and specific challenges. Programs range from intensive single-day sessions to comprehensive multi-week implementations of 20-40 hours. We work with groups of 5-25 participants and adjust pacing based on your team's experience level.

Offsite Intensive

Multi-day bootcamp at premium Mediterranean locations. Maximum 24 participants for focused, high-touch learning. All-inclusive program covering accommodation, meals, and team activities alongside training. Five days of immersive learning away from daily distractions. Participants leave with certification and internal playbooks ready for immediate implementation.

Online Sessions

Remote training for distributed teams across multiple time zones. Same methodology and hands-on approach as in-person sessions. Flexible scheduling that works around your sprint cycles and release schedules. Ideal for organizations with engineering teams spread across multiple locations.

Ongoing Support

For teams that want continued guidance, we provide access to our experts as you implement what you have learned. Regular check-ins, async support channels, and on-demand sessions for emerging challenges. Many teams find this valuable during the first 90 days of applying the framework at scale.

About Us

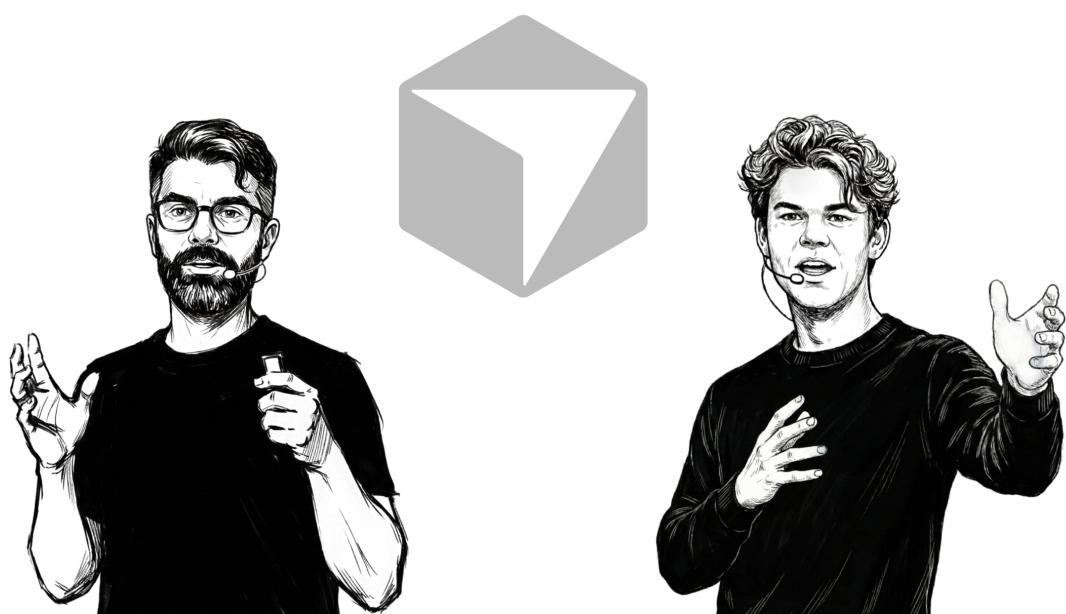
Cursor Workshop helps engineering organizations operationalize AI-assisted development. We are an Official Cursor Ambassador organization recognized by Anysphere. We train teams at Fortune 100 companies, DAX 40 enterprises, and high-growth technology organizations, including teams operating in regulated and security-sensitive environments.

Vasilis Tsolis, Co-founder

Official Cursor Ambassador for Greece. Founder of Cognitiv+ and senior enterprise software leader with a track record in B2B legal technology across the UK and Greece.

Rogier Muller, Co-founder

Official Cursor Ambassador for the Netherlands. CTO at a leading FinTech in the Netherlands. Founder of delta0, which acquires and operates enterprise software across compliance, legal and finance.



Next Steps

Organizations that operationalize AI-assisted development early gain a durable velocity advantage.

We work with a limited number of teams each quarter to maintain hands-on implementation quality. If your organization is already deploying AI coding tools but lacks standardized workflows, we should talk.

In this session we will:

- Assess your current AI adoption maturity
- Identify quality, security, and workflow risks
- Define the fastest path to measurable engineering productivity gains
- Recommend the training format and rollout plan that fits your team

Book a Strategy Session:

cursorworkshop.com/contact